
Commuter classification and behavior clustering: Beijing use case

Selene Baez Santamaria
s.baezsantamaria@student.vu.nl

Abstract

Public transportation, centered on subway and bus networks, is an data-rich domain that can benefit from data mining and machine learning techniques. The classification of commuters versus non-commuters/occasional travelers can help government, transport management and operators to better target their policies in order to improve the transportation network in large cities. Furthermore, characterizing commuters by behavior clustering can bring deeper insight into their needs and routines as a whole. This project proposes the usage of ensemble models for classification and clustering of public transport users. For this purpose, transit card data will be used, available from the city of Beijing, China.

Contents

1	Introduction	4
1.1	Transportation domain	4
1.1.1	Who are the commuters?	5
1.2	The city of Beijing	5
1.3	Motivation	6
1.3.1	Societal context	6
1.3.2	Scientific context	6
1.4	Thesis organization	7
2	Literature review	8
2.1	Data mining on transit card data	8
2.1.1	Volume of data	9
2.2	Classifying and clustering spatio-temporal data	9
2.2.1	Features	9
2.2.2	Classifying algorithms	9
2.2.3	Clustering algorithms	10
3	Research objective	12
3.1	Research questions	12
3.1.1	Definition of terms	12
3.2	Scope and structure	12
4	Methodology	14
4.1	The data	14
4.1.1	Special considerations	15
4.1.2	Labeled and unlabeled data	15
4.2	Data preprocessing	16
4.2.1	Cleaning	16
4.2.2	Trip parsing	17
4.2.3	Data patching	18
4.2.4	Transformation and Standardization	19
4.3	Data mining techniques	21
4.3.1	Spatio-temporal representation	21
4.3.2	Feature engineering	22
4.3.3	Ensemble models	23
4.3.4	Dimensionality reduction	23
4.4	Correlation analysis	23
5	Commuters identification	24

5.1	Features	24
5.2	Hypothesis	24
5.3	Model	24
5.4	Experiments	24
5.4.1	2 dimensional representation	24
5.4.2	3 dimensional representation	26
5.5	Results	26
6	Variable evaluation	27
6.1	Hypothesis	27
6.2	Qualitative	27
6.2.1	Interview	27
6.3	Quantitative	27
6.4	Feature selection	27
7	Commuters clustering	28
7.1	Feature engineering	28
7.1.1	Convolutional filters	28
7.1.2	Neural networks	28
7.2	Model	28
7.3	Experiments	28
7.4	Results	28
7.5	Expert judgment	28
8	Conclusion	29
9	Future work	30


1 Introduction

1.1 Transportation domain


Urban public transportation includes systems that are available for use by anyone in urban areas. Its facilities are commonly composed by buses, subway/metro lines, light rails, tramways, trains, taxis, and others. As a network, they provide service for the majority of citizens in urban areas.[22]

Figure 1 shows the passenger transport usage, as million passengers per kilometer. This represents the transport of a passenger for one kilometer. From the top image we note that United States, China, Germany, France, Italy, and United Kingdom constitute the six countries with the most passenger transport, according to their reported data from 2015 or later.[14]

Furthermore, historical data in the bottom image reveals the 15 years behavior for each of the aforementioned countries. Most of the countries show stability, with increase or decrease of less than .10 million passengers for European countries, and .5 million passengers for United States. China, however, shows a trend with steep increase for most of the selected years. In fact, comparing to its less than 1.2 million passengers in 2000, China doubled its public transport usage to 2.4 million passengers in 2015.



(a) Passenger transport data, per country.



(b) Historical data for the top six countries with most passenger transport usage.

Figure 1: OECD countries and their passenger transportation data.

Though it is a more sustainable alternative compared to private car usage, public transport usage has a significant environmental impact, affecting noise and air pollution. Diesel buses, which generally make up a major part of public buses, have large fuel consumption needs and contribute significantly to CO₂ emissions. Even eco-friendly alternatives such as hybrid diesel buses are sensitive to operating conditions, as their fuel consumption may increase by up to 50% when the on-board air conditioning is on.[26]

Consequently, public transportation directly relates to energetic demand, since its facilities are mostly petroleum or electrical based. In terms of global energy consumption, passenger transportation accounts for about 25% of the total world energy consumption. Furthermore, the transportation sector consumption increases at an annual average rate of 1.4.% [8] This may bring further economical implications for countries with high public transportation demand.

1.1.1 Who are the commuters?

A major proportion of public transport users is represented by commuters. These are regular users of public transit, with consistent spatio-temporal patterns in their travels. Driven by a routine, commuters travel back and forth from specific places, commonly represented by their home, work, school, or other similar locations.

As commuters are frequent users of public transit, the conditions of the public network directly influence their personal well being and generally impacts their quality of life. Intuitively, if the commuting experience is unpleasant, daily travel can bring distress to commuters and even repel them from using the public transport at all. Several studies have looked into public transit evaluation from different perspectives, including commuters' needs [12]. The most common aspects of it include: travel time, average speed, delays, accessibility, service coverage, crowded level, facilities quality, and fare rate. Weng et al [24] identified five categories (Convenience, rapid, Reliability and Comfort) that summarize commuters priorities when choosing to travel by public transport.

From both of the above, the large presence of commuters and their known needs and preferences, it follows that identifying commuters can help in creating a sustainable public transportation network. Public transit stakeholders should be able to understand the commuters' demands and its dynamics, consequently bringing long term planning and policies for improving the overall commuting experience.

1.2 The city of Beijing

The city of Beijing presents a special case of urbanization and rapid industrialization. This is reflected in a sudden population growth of 20% per decade since 1960, with the largest increase of 44% in the last ten years. The latest official census in 2010 reported the urban agglomeration of Beijing (including Beijing itself and its adjacent suburban areas) having a population of 19,612,368 people. The UN World Urbanization Prospects estimates the 2017 population at over 22 million inhabitants. [19]

As a result of the population explosion, many environmental and social resources are under pressure. From the environmental side, one of the most notable issues is related to air pollution, due to the significantly high pollutant emissions in the city [25]. Similarly, the city's downstream river pollution is serious, with most regions of the Yellow river being unable to comply with the lowest water quality standards. [23]

In the aspect of social resources, one of the main complications is the one of mobility. In Beijing, public transport is the dominant mode of transportation, accounted for 44.0% of all trips compared to 32.6% attributed to private cars [12]. In 2008, the total ridership was 6.5 billion travels. Though it is continually expanding, it is a fact that public transport is overcrowded, constantly reaching over 100% capacity [15].

Beijing public transport is composed of buses, subway and bicycles. The three types can be accessed by using a single smart card.

Bus: In 2015, there were 876 bus lines with 23,287 buses in operation. The bus network is the most extensive mode of transportation, expanding over 20,186 km. It observes an average daily traffic volume of 10.98 million passengers, with the highest daily volume reaching 13.07 million on one day. [3]

Subway: The Beijing subway has 18 lines with 334 stations, of which 53 are transfer stations. In 2015 it had an operating length of 554 km, with 5,024 vehicles running. [3] Its network is split by two operators: the state-owned Beijing Mass Transit Railway Operation Corp (operating 15 lines), and the joint Hong Kong venture Beijing MTR Corp (operating 3 lines). Beijing's subway has an average daily traffic volume of 9.11 million passengers, with a maximum recorded volume of 11.66 million passengers. As such, it is the second busiest metro system in the world, providing 3,410 million annual journeys. Compared to the service provided in 2012, the system observed a 39% increase in usage by 2014. It is also the second longest metro network, surpassed by Shanghai by only 21 km. [21]

Bicycles: Beijing first implemented public bicycle systems in 2012. As of 2015, in total, 67,000 bikes are available for rental with 2,700 pick up/drop off points spread across the city. [3]

1.3 Motivation

This project performs an interdisciplinary study between the areas of Artificial Intelligence and Metropolitan Transportation. It is focused on introducing data mining techniques to a data rich domain.

The area of Artificial Intelligence is able to provide dozens of prediction algorithms. Though constantly under refinement, it is time for state-of-the-art techniques to be applied to real and large impact situations to test their ability to deal with noisy streams. Comparably, given the ever growing complexity of urban mobility, domain experts must focus on analyzing trends and insights instead of curating and making sense out of raw data. Therefore, both areas benefit from this project.

1.3.1 Societal context

Looking at the social context of Beijing, there are several reasons justifying an in depth study of commuters.

First, a survey conducted in 2009 showed that 80% of the public transport passengers' complaints in Beijing were related to the network being slow and time-consuming, inconvenient to transfer, unpunctual and unreliable. Commuters use the public transport network regularly to perform their routines, and thus need efficient and reliable means of transportation. Flaws and failures in the network reduce the attraction of public transport. [15]

Second, the city of Beijing faces a large imbalance between residential and working areas. Due to urban expansion, most residents have been forced to move to suburban areas due to the lack of affordable housing [27]. Targeting this group brings the largest benefits to the public.

Third, government, transport management and operators can gain invaluable spatial and temporal insights regarding commuters' behaviors. This insight can lead to tangible results, including policies for increasing the efficiency of the public transit network, adjustable travel fares tailored to most relevant commuters' patterns, incentives to relieve peak hours and thus traffic congestion, urban planning for residential and industrial land use, and others.

1.3.2 Scientific context

From a scientific point of view, mobility patterns in metropolitan areas follow swarm behaviors. Based on individual travels and routines, travelers exhibit distinguishable characteristics on a larger scale. Both levels of understanding are crucial for Transportation experts.

Surveys are the most widely used research method in metropolitan transportation studies. They are targeted to reach travelers on an individual level, while aggregated measurements are taken to investigate their collective behavior. The analysis is usually based on statistical methodology.

These methods have several disadvantages. On the one hand, surveys are costly to implement, and in general have problems related to small populations and non-representative samples. Even when the latter problems are escaped, the typical quality versus quantity trade off is present, reducing the confidence of the collected information. Furthermore, surveys are based on self-report, which by itself has bias problems. On the other hand, aggregated measurements miss the interactions between individuals that cause the collective behavior.

Big Data In the last years, smart card systems have become more popular, making it possible to monitor travelers transactions and facilitating fare collection. Several cities have implemented such systems, for example the Octopus card in Hong Kong[5], Oyster card in London [2], OV-chipkaart in The Netherlands [6], and Yikatong card in Beijing [4], to name a few.

In Beijing, over 90% of public transit users are smart card holders. There is a significant incentive for using the Yikatong smart card since bus rides are heavily subsidized (the user has only to pay 50% of the full price)[10]. Moreover, the Yikatong smart card system is also integrated with taxi, electricity and sewage payments.

Clearly, technology for transit data collection is broadly spread around the world. Transit smart cards are capable of recording spatio-temporal information at an individual level over long periods of time. This generates a large volume of historical data that only big data techniques can deal with. Furthermore, machine learning and other mining methods specialize in analyzing disaggregated complex information. As such, bringing state-of-the-art techniques into the Metropolitan Transportation domain can aid to unravel massive human behaviors and reveal patterns and trends in mobility.

1.4 Thesis organization

This Thesis is organized as follows:

First we perform a literature review to look for previous work on mining smart card transit data and for specific state-of-the art methodologies for classification and clustering complex spatio-temporal data. Subsequently, we establish the scope and research objectives of this project. We continue to describe the methodology thoroughly, including the data and the approach to mine it. Following this description, we identify three distinct stages of the project (Commuters classification, Variables evaluation, and Commuters clustering) and report their corresponding experimentation. Next, we discuss the findings and gather conclusions. Finally, future work opportunities are explored.¹

¹more detailed

2 Literature review

In this section we look at studies within the last decade that are related to smart card transit data, and the clustering techniques for their mining.

2.1 Data mining on transit card data

With the introduction of smart card systems in large cities, several studies have aimed to extract knowledge from the large amounts of data collected. Though different in their methodology, results concerning commuters are duplicated across studies. The spatial and temporal regularity of commuters' travel behavior is evident in their smart card data.

Morency et al. study spatio-temporal variability in Canadian smart card data. On the one hand, they examine spatial variability by measuring the number of distinct stops a smart card user visits, and the frequency of each stop. On the other hand, they examine temporal variability by clustering the boarding times of each type of smart card. Using these features, they observe the week to week variability for each of the five types of transit card available (Adult-interzone, Adult-express, Adult-regular, Elderly and Student). Their findings show that commuter types of cards visit a smaller range of bus stops compared to non-commuter types. Therefore, a small number of stops account for a high proportion of commuter's boardings. Additionally, commuters have the highest proportion of zero-boarding days on weekends [13].

Bhaskar et al. are concerned with passenger segmentation using Australian smart card data. First, they perform a two level DBSCAN algorithm for investigating spatial patterns, where the first level clusters Destination stops and the second level clusters Origin stops. From this they extract frequent Origin-Destination (O-D) pairs. Separately, they applied DBSCAN to temporal features to determine most frequent boarding times. As such, they characterize each user by the percentage of journeys they perform between the regular O-D, and the percentage of journeys they perform during their habitual times. Users with at least 50% spatial and temporal regularity are thus classified as transit commuters; while users with no evident spatial or temporal pattern are classified as irregular passengers. The authors find that while most (64%) of the passengers riding the public transit are irregular passengers, it is transit commuters who bring the most (46%) revenue. Furthermore, they find that irregular passengers prefer high frequency routes significantly more than transit commuters, arguing that commuters are usually on a time habit, and thus are more willing to check and adapt to public transit timetables. [1]

Tu et al. follow a supervised learning approach to classify public transit users in Beijing as commuters or non-commuters. In order to produce labeled data, they convey an online survey asking for travel patterns and smart card ID. Matching the ID to the journeys recorded by smart card during the span of one week, they collect records associated to 978 travelers. The classification is then performed by a Support Vector Machine (SVM), which reaches up to 94.24% accuracy. [20]

Langlois et al. present an innovative representation for smart card data. Using four weeks worth of data from London Oyster cards, they represent the card information as a time-ordered sequence of inferred activities. 11 clusters are found and characterized by evaluating socio-demographic variables like age, employment, annual household income, children per household and vehicles per household. The authors further grouped the clusters under "working day", "home bound", "complex activity pattern" and "interrupted pattern" categories. Their findings show that four clusters, grouped under the "working day" category have significantly different activities during weekdays as compared to weekends, with some avoiding transit during the weekends and others visiting different areas. Four more clusters, grouped under the "home bound" category, are characterized by staying mostly at their primary area and low number of traveled days. [9]

One of the latest work on the field corresponds to Ma et Al. The objective of their work is to determine a scoring function for travelers that can correctly identify them as commuters, or non-commuters. In their work, they cluster stops using an improved DBSCAN algorithm. They engineer features for representing the frequency in which travelers follow spatio-temporal patterns. Travelers are then clustered according to these features following the ISODATA algorithm. As an output of the clustering, optimal cutoff levels in the scoring function were determined. As a result, evaluating a traveler does not depend on clustering centroids, but only on calculating the commuting score. This,

as expressed by the authors, reduces computing time and treats each traveler independently from the others, which is not true for clustering algorithms [10].

A common practice, as used by [10], [9], and [13] is to divide the day into -hourly or half-and-hour-time bins. Bhaskar et al. recognize this as a problem in the field, by pointing out that this design choice segregates journeys from 9:59 AM and 10:01 AM even though they intuitively belong to the same behavior.

2.1.1 Volume of data

The volume of data collected by smart card systems is massive and is usually impossible to analyze all of it at once. The volume of the samples analyzed by previous work ranges from hundreds of smart cards to tens of millions of smart cards, leading to up to hundreds of millions of individual smart card transactions. The details are summarized in Table 1.

Authors	Year of publication	Records	Unique smart cards	Time span
Tu et al. [20]	2016	8,067	978	one week
Morency et al. [13]	2007	2.2 million	7,118	277 days
Langlois et al. [9]	2016	3 million	33,026	four weeks
Bhaskar et al. [1]	2015	34.8 million	1 million	4 months
Ma et al. [11]	2013	Unknown	3 million	one week
Ortega [16]	2013	65 million	5.7 million	one week
Ma et al. [10]	2017	364 million	18 million	one month

Table 1: Volume of data analyzed by different authors

2.2 Classifying and clustering spatio-temporal data

2.2.1 Features

Human mobility is intrinsically tied to spatio-temporal properties. Still, the greatest amount of studies analyze public transit journeys by separating spatial features from temporal features. In general, scalar aggregated features are used for users characterization. Some examples are:

- **Frequency indicators:** number of traveled days [1] [9] [10], number of journeys [1], number of times a stop was visited [13], number of days with zero boardings [13], most frequent home/work stop [10], most frequent home/work route [10], most frequent departure time from home/work [10], number of trips to the most frequent home/work stop [10], number of trips following the most frequent home/work route [10], number of trips during most frequent departure time from home/work [10]
- **Range/coverage indicators:** distinct stops visited [13], spread of days between the first and last journey [9]
- **Calendar-based indicators:** observed day [13], day of week [13]

Different from all of the above, Langlois et al. follow a unique methodology for engineering features. First, they represent the travel data per user using a three dimensional matrix where x represents the day in the four week period, y represents the hourly time bin, and z represents the area where the inferred activity took place, encoded as a one hot vector. The authors perform Principal Component Analysis (PCA) for dimensionality reduction, based on Eagle and Pentland’s eigenbehaviours [7]. An analysis of the average correlation of the first 13 components, results in the selection of the first 8 components as the most informative and stable. The projections of a user sequence onto these components (called weights) constitute the features to be clustered using k-means. [9]

2.2.2 Classifying algorithms

The domain of Metropolitan transportation faces a specific problem: although smart card systems have allowed massive collection of data, this data is not labeled regarding commuting behaviors. Additionally, obtaining labels for smart card data is expensive and unreliable, since it has to be acquired through surveys or interviews. Furthermore, even when labels are obtained, the amount of

labels obtained is often insufficient for big data analysis. It is due to these reasons, that most studies are inclined to use unsupervised learning techniques.

One of the few studies that uses labeled data corresponds to Tu et al. They obtain 978 labeled records, with an almost equal distribution of records over both classes (49.18% related to commuter samples and 51.82% related to non-commuter samples). They solve the issue of limited samples by selecting a model that is not heavily affected by sample size: Support Vector Machines. Their results report a 94.24% accuracy over a test set of 295 samples.

2.2.3 Clustering algorithms

If labeled data is not available, then unsupervised learning techniques must be applied. There is a large variety of clustering algorithms available nowadays, however not all of them are suitable for all types of data and purposes.

Hierarchical clustering

Langlois et al. use agglomerative hierarchical clustering for areas clustering. In order to infer the user-specific activities, all stops or stations visited by each user are clustered by merging the two closest areas until a threshold distance is reached. Their algorithm also considers the distance between stops and the frequency of travel between them. Therefore, different activities are likely to be associated with different areas [9].

Partitional clustering

K-means algorithm is the most widely used method for partitional clustering. It requires having a predefined number of clusters to fit the data to.

Morency et al. use K-means for clustering hourly boarding times according to card type. They apply Hamming distance (representing the percentage of data between two elements) and a combination of batch and online updates. Through empirical tuning, they select to find four clusters per card type. It is worth noting that by using a card-day unit, they allow a card to belong to a different cluster according to the day of travel. As every card type is composed of four boarding patterns, travelers are not restricted to follow a routine everyday, but can exhibit different behaviors on different days. For example, the Adult-regular card type contains a 9:00AM-and-5:00PM-boarding cluster and a no-boarding cluster. Thus, a user of this card could belong to the first cluster on weekdays and to the second cluster on weekends. [13]

Bhaskar et al. apply K-means for binary classification purposes. As such, they classify frequent and infrequent transit users, using the number of traveled days and the number of journeys made as features. Unfortunately, K-means performs poorly since no distinct clusters are evident. The most likely cause for the previous is the strong correlation between traveled days and journeys, combined with the authors oversight of whitening and standardization techniques. [1]

Langlois et al. use K-means to find clusters of activity sequences. They employ specialized sampling techniques, like bootstrapping, to deal with big data. Moreover, they tune the algorithm parameters using the DB-index, which is the ratio of the within cluster distances to the across cluster distances. They find two optimal number of clusters (4 and 11), out of which they select the largest to provide the most detailed segmentation. They further perfection the algorithm by using k-means++ initialization over 150 replications. Additionally, this paper acknowledges that clustering techniques are sampled based, which means different samples may find different optimal solutions. The authors validate their approach by analyzing the stability of the clusters over samples obtained at different points in time. By extracting the same number of clusters and fitting the samples to each set, they find that 91% of users are assigned to their equivalent clusters. [9]

Density based clustering

Density based algorithms excel at dealing with anomalies, since they ignore low density areas and interpret them as noise. They do not require a redefined number of clusters and adapt to find clusters of any size. The required parameters for DBSCAN are a maximum reach distance ϵ and the minimum number of points per cluster.

Bhaskar et al. use three DBSCAN algorithms to cluster Origin stops, Destination stops, and boarding times. For each of the previous, they tune the algorithm parameters by fixing a domain reasonable ϵ (1000 m walking distance or 5 min variance in boarding time), and selecting the minimum points by

comparing the percentage of data considered to belong to any cluster as opposed to data considered to be noise given the par-specific parameters [1].

Ma et al. use an improved DBSCAN algorithm to cluster bus/subway stops. In their approach, abnormal stops are not considered noise, but are allowed to be re-clustered by splitting large clusters into several smaller clusters.

Though clustering algorithms are common in the field, they are not always used for classifying users. For example, Bhaskar et al. use density based clustering for engineering regularity features. However, the classification of users is rule-based according to which feature (spatial or temporal regularity) is stronger in each user [1]. Morency et al. use partitioning clustering to characterize existing user categories according to their boarding times [13].

As a conclusion, we note that while there has been research applying basic clustering and classification algorithms, most studies lack further specialized data mining techniques for preprocessing data, tuning algorithms parameters, and/or visualizing results.

3 Research objective

The objective of this project is to identify and characterize commuters in the city of Beijing by using smart card data. As such, the underlying goal is to find patterns in the spatio-temporal data of public transport travelers, and interpret the results using domain knowledge by experts in the Transportation domain.

3.1 Research questions

The main objective is further broken down into answering the following research questions:

1. Given the same dataset, how does an ensemble model compare to the SVM model reported by Tu et al [20]?
2. Given an extended dataset, how accurately can commuters and non-commuters be identified using an ensemble model?
3. What is the minimal set of features needed to reach an acceptable (at least 95%) accuracy in classification?
4. How can spatio-temporal features be analyzed as a unit?
5. To what extent is clustering commuters by its behavior informative to transportation specialists?

3.1.1 Definition of terms

A commuter is a public transit user whose smart card data reveals repeatable patterns in time and space. Though commuters are usually associated with Monday to Friday 9:00am to 5:00pm schedules, in this work we extend the definition to any routine travel pattern. This flexibility allows us to include travelers with stable yet rare commuting schedules, such as night workers, weekend workers and evening workers.

A trip is a sequence of smart card transactions, including transfers, performed by the same user to travel from an origin to a destination. A trip is also represented as a record in the data, as it will be further explained in Section 4.1

A transfer is a change in transportation mode, or a change in vehicles whenever a smart card has to be checked within the same transportation mode. Transportation modes include Bus, Subway, and Bike.

We make the assumption that smart card IDs and users have a one to one relationship, meaning each user has exactly one card and each card is used by exactly one user. As discussed with domain expert Quian Tu, although some people may own more than one card, this is a minority. Thus, the assumption holds for the majority of travelers.

3.2 Scope and structure

This project is divided three main stages:

PART I Classify commuters versus non-commuters by using an ensemble model. In this part we focus on research questions 1 and 2. We examine the performance of ensemble models given different datasets, and make a comparison with Tu’s SVM model [20]. This part corresponds to Section 5.

PART II Revise the feature space in order to identify the most informative features and disregard redundant information. In this part we focus on research question 3. An extensive analysis of spatio-temporal properties will be done, combining transportation domain knowledge and statistical tools. This part corresponds to Section 6

PART III Commuters clustering according to patterns in their travel behaviors. In this part we focus on research questions 4 and 5. Different clustering techniques will be combined via consensus clustering. The final clusters will be analyzed and interpreted to find distinctive characteristics that may be deemed as informative for transportation specialists. This part corresponds to Section 7

²make more technical
³make flowchart of processes in project

4 Methodology

4.1 The data

Every record in the data represents a trip performed by a specific smart card. As such, it contains the following data fields:

- Data date: Year, month and day that the trip was made
- Card code: Card identification number
- Path link: Mode of transportation. B stands for bus, R for subway, Y for bicycle. Transfers between modes are shown by a dash.⁴
- Travel time: Time spent in vehicles, measured in milliseconds
- Travel distance: Distance traveled, measured in meters as performed by route.
- Transfer number: Number of changes in travel mode during the trip.
- Transfer total time: Total time spent in transfer, measured in milliseconds
- Transfer average time: Time spent in transfer, divided by number of transfers. Measured in milliseconds
- Start/End time: Time stamp of when the trip started/ended. Date and time up to milliseconds precision
- On/Off small traffic area: Integer ranging from 1 to 1911
- On/Off middle traffic area: Integer ranging from 1 to 389
- On/Off big traffic area: Integer ranging from 1 to 60
- On/Off ring road: Integer ranging from 1 to 6
- On/Off area: Integer ranging from 1 to 18
- ID: record identification number created by joining the following: hour of the beginning of the trip | time stamp of beginning of the trip | card code performing the trip
- Transfer detail: Mode of transportation, as well as line/route number and stations for boarding and alighting. More detail provided in Section 4.2.2

Full privacy of card users is ensured, as there is no personal data linking card codes to specific individuals.

The traffic zones (small, middle and big areas) are divided by the Beijing Municipal Institute of City Planning and Design (BICP). They are specific in different degrees, as shown in Figure 2. In general, the division principles correspond to the geopolitical environment and administrative planning, for example roads, villages and others. The 6 ring road and 18 areas districts are divided by the Beijing Municipal Government. The division is unique in Beijing. The 18 districts and counties are shown in Figure 3. According to domain expert PhD. Liang Quan, these divisions are sufficiently informative for traffic analysis [18].




Figure 2: Traffic zone division

Every day, more than 13 million records are collected, with approximately 5 million corresponding to subway trips, 8 million corresponding to bus trips and 100,000 corresponding to bicycle trips.

⁴Example: B-B represents a Bus to Bus transfer.

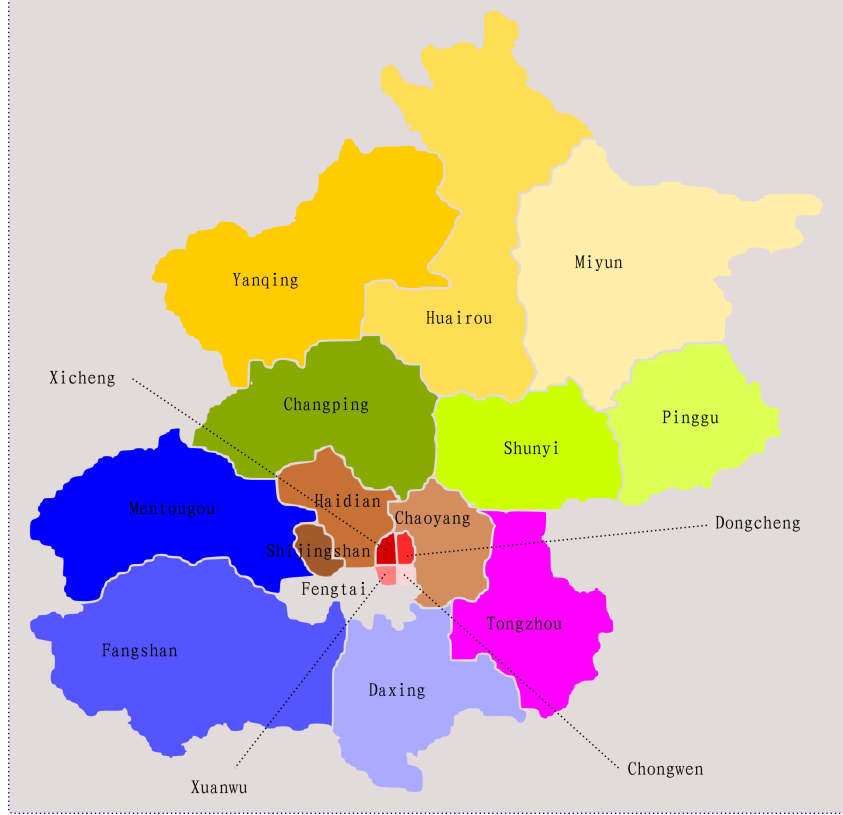


Figure 3: Beijing's Districts and its Counties

4.1.1 Special considerations

The previous description corresponds to the data as delivered by the Beijing Transportation Research Centre. As such, it is the result from processing the raw records at the collecting phase. Some special considerations concerning this processing are explained below:

Travel distance by bike: Since bicycles do not have predefined routes, the distance cannot be directly recorded. However, it is inferred by using the travel time and a static average speed for cyclists.

Subway transfer: Transfers between subway lines of the same operator cannot be tracked since a single check-in gives access to all the subway network. In order to infer the transfer detail, the A* algorithm is used to calculate the most likely transfer sequence, given the boarding and alighting stations. Similarly to the bicycle missing information, the transfer time inside the subway system cannot be directly recorded. Using a static average walking speed and the known distance in transfer stations, the transfer time is calculated.

Transfer information: The path link and transfer number fields are extracted from the transfer detail field. Similarly, the transfer average time is calculated from the transfer total time and transfer number fields.

4.1.2 Labeled and unlabeled data

Commuter classification: Classification is a supervised learning task, where every training data sample requires an associated label determining its true class. In case of commuter classification, this translates to having smart card codes associated with either a "commuter" or "non-commuter" label. Such data is expensive and limited since it can only be obtained by asking the users directly if they are commuters or not. Thus, in general, annotated data is not available, and labeling new records falls beyond the scope of this project.

As a solution for the above, we take advantage of the dataset used by Tu[20]. This dataset corresponds to trip records performed during a week in January 2015, and it contains labels for 978 smart cards, collected and validated via surveys. The original dataset distribution is composed by:

- 6439 records of 481 commuters
- 1628 records of 497 non-commuters

For this project, the Beijing Transportation Research Centre has provided us with one month worth of data, corresponding to January 2015. In order to construct an extended labeled dataset, we take these 978 labeled smart card IDs and search for their corresponding records in the one month sample. This dataset is used for Part I (Section 5) and Part II (Section 6) of this project.

Commuter clustering: In order to further cluster commuters, 100,000 smart card codes are sampled from data corresponding to November 2016. The month of November is chosen because it does not overlap with holidays and has a relatively stable weather thus diminishing the variance between bicycle and bus/subway traveler preferences. The year is chosen to reflect a more recent characterization of travelers.

Using the commuter classifier from Part I, we filter to keep only commuter related records. This dataset is used in Part III (Section 7) of this project.

4.2 Data preprocessing

4.2.1 Cleaning

As first step for preprocessing the data, we eliminate faulty records. The different filters are:

1. Eliminate records with missing data: 10.93% records eliminated
2. Eliminate records with missing travel details: 1.58% records eliminated
3. Eliminate records with travel time ≤ 0 : <0.01% records eliminated
4. Eliminate records with travel distance ≤ 0 : 9.82% records eliminated
5. Eliminate records linked to users with insufficient trips: 25.89% records eliminated when requiring at least 2 trips.

The first four filters aim to eliminate records with missing fields, which already reduces the dataset to 77.66% of its original size. The last filter aims to eliminate users with insufficient number of trips. These users represent straightforward cases of irregular public transit users are unlikely to be commuters. Therefore, the minimal number of trips per user is a parameter that can be tuned.




Figure 4: Number of trips distribution. 2500 record sample.

Figure 4 shows the distribution of number of days in a single day. It shows that most people perform two trips per day. Fixing the minimal number of trips to 60, which is equivalent to an average of two trips per day, the final dataset contains 51.76% records available for usage.

4.2.2 Trip parsing

The trip details obtained from the records are given in Chinese, with descriptors containing a combination of numbers and text. In order to extract boarding/alighting route features, the descriptors must be parsed.

We parse the trip details using a combination of two techniques: regular expressions and tokenization.

Regular expression Since a trip may include transfers, we define a trip to be composed of one or many rides. Each ride is carried out in a single travel mode.

In order to obtain the elements of each ride we look at the pattern per travel mode.

$$\begin{aligned}
 BIKE &= (bike.STOP - STOP) \\
 SUBWAY &= (subway.LINE : STOP - LINE : STOP) \\
 BUS &= (bus.ROUTE(DIRECTION - DIRECTION) : STOP \\
 &\quad - ROUTE(DIRECTION - DIRECTION) : STOP)
 \end{aligned}$$

where the upper-case text corresponds to placeholders for ride elements, the lower-case text corresponds to the English translation of the descriptor in Chinese, and the punctuation (parentheses, dots, colons and dashes) correspond to separators between ride elements.

Unifying the mode-specific patterns, we describe a ride and a trip using regular expressions:

$$\begin{aligned}
 RIDE &= (MODE.[LINE/ROUTE:]?STOP - [LINE/ROUTE:]?STOP) \\
 TRIP &= RIDE[- > RIDE]?
 \end{aligned}$$

where elements surrounded by squared brackets and followed by a question mark (e.g. $[ELEMENT]?$) correspond to optional elements. We note that when parsing bus details, we

disregard the route direction. This decision is motivated to fit both subway lines and bus routes to a single pattern, noting that the direction of the route does not affect the path of the route itself.

Tokenization Once the elements of a trip are extracted, they must be substituted with numerical IDs. These IDs are not available from the Beijing Institute of Transportation, thus three different vocabularies are created for subway lines, bus routes and combined stops correspondingly.

Usually, bus routes are identified by a number. However, in Beijing a single bus route number can be associated to different paths. Such is the case of night, express and special cases of a bus route, which follow different paths even if they are described with the same number. For this reason we create a vocabulary with all unique parsed routes according to their full description and not only their number.

Examples of cleaned routes are shown in Figure 5

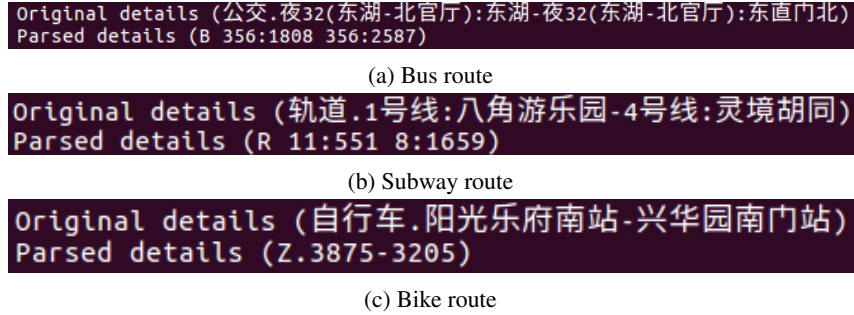


Figure 5: Examples for parsed and tokenized trip details.

4.2.3 Data patching

We note that the number of transfers and the path link fields of some records do not correspond to the information in their trip details. According to domain expert PhD. Tu Qiang, this must be recalculated [17]. Figure 6 shows the distribution of the number of transfers per trip before and after patching. ⁵




Figure 6: Transfer number distribution before and after recalculation.

⁵piechart might be better?

Our distribution shows that most trips are performed without transfers, which is consistent with other studies findings [1].

4.2.4 Transformation and Standardization

In data mining, it is a standard practice to perform whitening. This technique eliminates correlations between features, which is desirable in most cases. However, the domain of Metropolitan Transportation some of these correlations are highly important, and should not be discarded. This is the case of total travel time and distance, as shown in Figure 7. For this reason, we choose to only standardize the features and keep the correlations.




Figure 7: Travel distance vs travel time. 2500 record sample.

Travel time, travel distance, total transfer time and average transfer time were standardized by subtracting the mean of each distribution and forcing a unit standard deviation.




Figure 8: Time distribution before and after preprocessing. 2500 records sample.




Figure 9: Distance distribution before and after preprocessing. 2500 records sample.

Figure 8 and 9 show travel time and distance follow a truncated Gaussian distribution $\mathcal{N}(\mu = 1, \sigma^2 = 1)$. Since the nature of the data prevents negative values (time and distance must be positive), the original distribution is truncated at 0. Standardization maintains the shape of the distribution, but shifts and contracts it to be closer to zero values.

The mean travel time is 55 minutes, and the mean travel distance is 10 kilometers.

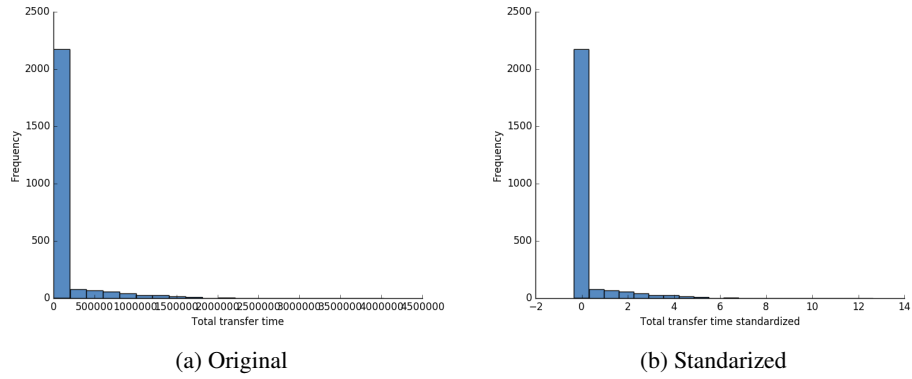


Figure 10: Total transfer time distribution before and after preprocessing. 2500 records sample.




Figure 11: Average transfer time distribution before and after preprocessing. 2500 records sample.

Figure 10 and 11 show that transfer times, both total and average, follow distributions with very long tails. As mentioned before, most trips are performed without transfers, which explains that most of the trips have transfer times equal to zero.⁶

Time bins

Regardless of its criticism, using hourly time bins is standard practice in the field and has shown sufficient to examine temporal data [9] [10] [13]. Therefore, in this project we follow the same technique and extract only the hour of the start and end of each trip.




Figure 12: Distribution of start/end hours for trips. 2500 records sample.

From Figure 12, we note that our data follows the expected distribution for the domain, showing clear morning and evening peaks. Furthermore, we note that boarding and alighting patterns during the peaks hours are shifted by one hour. This is explained by the previous finding, that the mean travel time is almost one hour.

4.3 Data mining techniques

The work for this project is coded in Python, using libraries for data and numerical manipulation (*Pandas*, *numpy*), regular expressions (*re*), machine learning (*sklearn*), neural networks(*Theano*).

4.3.1 Spatio-temporal representation

In this work we propose a 3 dimensional data representation to contain the monthly travel information of a user. This is shown in Figure 13.

⁶do not include transfer 0. focus on tail behavior




Figure 13: Spatio-temporal data structure.

Inspired by [9], the x-y plane constructs a temporal structure between days of the month and hours of the day. The crucial advantage of this structure lies in its local properties. Similar to the case of image processing, in this representation a temporal pixel is simultaneously influenced by what happened in the previous/following hours (y axis), and on the previous/following days (x axis).

As for the z plane, each layer contains a trip feature. In Figure 13 boarding spatial features are portrayed in green, alighting spatial are portrayed in red, and other types of features (such as travel time, travel distance, transfer number, transfer total time, etc) are portrayed in yellow.

Therefore, each temporal pixel may contain a trip feature vector, which expands several layers deep. Considering that even regular public transport users do not perform more than 6 trips a day as shown by the number of trips distribution in Figure 4, the proposed representation is sparse, since only a few time pixels are populated with trips.

4.3.2 Feature engineering

The spatial features to be included in the representation are:

1. Small traffic area
2. Middle traffic area
3. Big traffic area
4. Ring road
5. District
6. Mode
7. Line/Route
8. Stop

These are repeated for boarding and alighting information.

The general trip features to be included in the representation are:

1. Travel time
2. Travel distance
3. Transfer number

4. Transfer total time
5. Transfer average time
6. Start hour
7. End hour
8. Number of trips

4.3.3 Ensemble models

Ensemble models benefit from combining non-correlated prediction methods. Weak classifiers might correct each other in specific hard cases. Ensemble models are chosen for this project because of its robustness and modularity. Starting from a few simple classifiers, assembled via aggregation methods, the model can grow larger or more complex as needed.

Supervised learning As proven by Tu [20], weak classifiers like an SVM prove to be sufficient to identify commuters. This hints to extend the ensemble model with other similar weak classifiers like decision trees, Bayesian classifiers or multilayer perceptron. Bagging will be used to ensemble their predictions.

Unsupervised learning The equivalent to ensemble models in unsupervised learning is consensus clustering.

4.3.4 Dimensionality reduction

Considering these are 24 features, we have $24 * 30 * 24 = 17,280$ temporal pixels per user. Given the high dimensionality and the sparsity of the structure, we will perform dimensionality reduction.

Taking advantage of the local properties of the proposed structure, we can apply convolutional filters⁷ to reduce the dimensionality to a more manageable number. The end result will be used as features for clustering commuters in Part III (Section 7) of this project.

4.4 Correlation analysis

A different method for reducing the number of features in the model is to perform a factor analysis. The independent variables may be evaluated against each other to test their correlation using the Chi-Square test. The features that are strongly correlated are redundant, and thus can be eliminated from the model.

⁷CNN chosen because of local properties, is PCA also local?

5 Commuters identification

5.1 Features

A first instance of the model will use features as used by Tu [20] for a fair model comparison. A second instance will contain all the variables described in Section 4.3.2

Tu features and representation

Proposed features and representation

5.2 Hypothesis

As suggested by Tu [20] results, the data is almost linearly separable thus simple classifiers such as decision trees may suffice.

The 3 dimensional representation will prevent overfitting.

5.3 Model

We compare a single SVM, ensembled random forests.

5.4 Experiments

5.4.1 2 dimensional representation

SVM

Accuracy: 94.66%




Figure 14: SVM confusion matrix.

Random Forest

Accuracy: 97.54%




Figure 15: Random forest confusion matrix.

AdaBoost

Accuracy: 98.24%

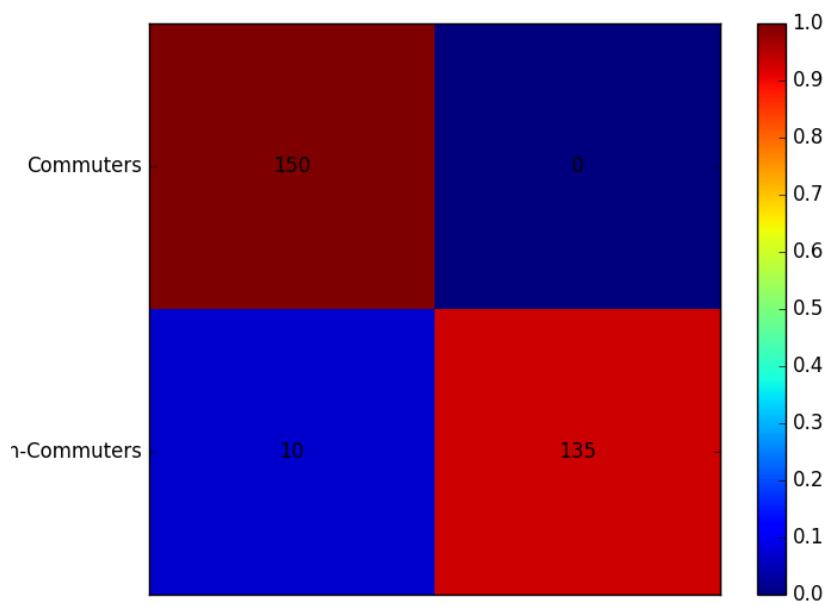


Figure 16: Random forest confusion matrix.

5.4.2 3 dimensional representation

5.5 Results

Confusion matrix

6 Variable evaluation

6.1 Hypothesis

One of the main focuses of the second phase of this thesis is to determine the appropriate level of detail in the area to be taken into account.

Small, middle and big areas are hierarchical divisions. Middle and small divisions have more precision but maybe are not needed. On the other hand, big area divisions might not capture the changes for people who live and work/study in the same district.

6.2 Qualitative

Exploration: Experts opinion

6.2.1 Interview

We interview Liang Quan as an Transportation domain expert.

- To what extent do people live and work on the same district?
- What level of detail in traffic areas do you think is appropriate?

6.3 Quantitative

Analysis: Correlation

6.4 Feature selection

7 Commuters clustering

7.1 Feature engineering

The features to be fed into the clustering algorithms are obtained by dimensionality reduction. As mentioned before, convolutional filters can exploit local information.

7.1.1 Convolutional filters

In order to examine weekly patterns, our the first layer of convolutional filters will have an x dimension of 7 with a stride of 7. Assuming that only the hour previous and after a trip affects the trip, the filter will have a y dimension of 3 with a stride of 1. We do not perform padding, since this would have significant implications in the travel behavior of the user.

Following this formula:

$$output = \frac{input - filter}{stride} + 1$$

We find that the output size after the first convolution is 4×22 . Considering 15 features, this reduces the dimensionality to $4 \times 22 \times 15 = 1320$

Goal is to have a $4 \times 3 \times 3 = 36$ structure

⁸

7.1.2 Neural networks

A neural network is constructed to apply the convolutional filter. After the convolution layer, ReLU is applied as activation function.

⁹

7.2 Model

The model will be created by combining results from k-means and DBSCAN algorithms. This technique is called consensus clustering.

First, a DBSCAN will be run to determine the appropriate number of clusters. Then, this number will be set as parameter for k-means. The results will be integrated using hard ensemble clustering, specifically the meta-clustering algorithm.

7.3 Experiments

7.4 Results

7.5 Expert judgment

⁸what about depth? can I have a filter that is 8 units deep, with 8 stride (features are divided as 8 spatial boarding, 8 spatial alighting and 8 transfer related)? How could this be applied?

⁹standard practice in deep learning, but not sure if it should be applied in this case

8 Conclusion

9 Future work

References

- [1] Ashish Bhaskar, Edward Chung, et al. Passenger segmentation using smart card data. *IEEE Transactions on intelligent transportation systems*, 16(3):1537–1548, 2015.
- [2] Philip T Blythe. Improving public transport ticketing through smart cards. In *Proceedings of the Institution of Civil Engineers, Municipal Engineer*, volume 157, pages 47–54. Citeseer, 2004.
- [3] Beijing Transportation Research Center. 2016 annual report on traffic development in beijing. http://www.bjtrc.org.cn/InfoCenter/NewsAttach/2016%E5%B9%B4%E5%8C%97%E4%BA%AC%E4%BA%A4%E9%80%9A%E5%8F%91%E5%B1%95%E5%B9%B4%E6%8A%A5_20161202124122244.pdf, 2016. Accessed on 21 April, 2017.
- [4] Mo Lim Chan. *Tactical implementation model for the smart card payment system for metro operation*. PhD thesis, City University of Hong Kong, 2010.
- [5] Patrick YK Chau and Simpson Poon. Octopus: an e-cash payment system success story. *Communications of the ACM*, 46(9):129–133, 2003.
- [6] Gerhard de Koning Gans. *Analysis of the MIFARE Classic used in the OV-chipkaart project*. PhD thesis, Master’s thesis, Radboud University Nijmegen, 2008.
- [7] Nathan Eagle and Alex Sandy Pentland. Eigenbehaviors: Identifying structure in routine. *Behavioral Ecology and Sociobiology*, 63(7):1057–1066, 2009.
- [8] US EIA. Energy information administration (2016), international energy outlook 2016, with projections to 2040. Technical report, DOE/EIA-0484, 2016.
- [9] Gabriel Goulet Langlois, Haris N Koutsopoulos, and Jinhua Zhao. Inferring patterns in the multi-week activity sequences of public transport users. *Transportation Research Part C: Emerging Technologies*, 64:1–16, 2016.
- [10] Xiaolei Ma, Congcong Liu, Huimin Wen, Yunpeng Wang, and Yao-Jan Wu. Understanding commuting patterns using transit smart card data. *Journal of Transport Geography*, 58:135–145, 2017.
- [11] Xiaolei Ma, Yao-Jan Wu, Yinhai Wang, Feng Chen, and Jianfeng Liu. Mining smart card data for transit riders’ travel patterns. *Transportation Research Part C: Emerging Technologies*, 36:1–12, 2013.
- [12] Zidan Mao, Dick Ettema, and Martin Dijst. Commuting trip satisfaction in beijing: Exploring the influence of multimodal behavior and modal flexibility. *Transportation Research Part A: Policy and Practice*, 94:592–603, 2016.
- [13] Catherine Morency, Martin Trepanier, and Bruno Agard. Measuring transit use variability with smart-card data. *Transport Policy*, 14(3):193–203, 2007.
- [14] OECD. Passenger transport (indicator). 10.1787/463da4d1-en, 2017. Accessed on 10 April, 2017.
- [15] Beijing Municipal Committee of Communications and Beijing Transportation Research Center. Research on beijing’s public transportation commuting transit network. https://www.esmap.org/sites/esmap.org/files/10282009102930_Beijing_Transport_finalReport.pdf, 2009. Accessed on 21 April, 2017.
- [16] Meisy Andrea Ortega-Tong. *Classification of London’s public transport users using smart card data*. PhD thesis, Massachusetts Institute of Technology, 2013.
- [17] Tu Qiang. personal communication.
- [18] Liang Quan. personal communication.
- [19] World Population Review. Beijing population. <http://worldpopulationreview.com/world-cities/beijing-population/>, 2016. Accessed on 21 April, 2017.

- [20] Qiang Tu, Jian-cheng Weng, Rong-Liang Yuan, and Peng-fei Lin. Impact analysis of public transport fare adjustment. *Traffic Engineering & Control*, 57(2), 2016.
- [21] UITP. World metro figures, statistics brief. http://www.uitp.org/sites/default/files/cck-focus-papers-files/UITP-Statistic%20Brief-Metro-A4-WEB_0.pdf, 2015. Accessed on 21 April, 2017.
- [22] Vukan R Vuchic. Urban public transportation systems and technology. 1900.
- [23] Kunyan Wang, Qiong Luo, and Xueying Zang. Studies on ecological environmental carrying capacity in beijing, tianjin, and hebei region. In *Report on Development of Beijing, Tianjin, and Hebei Province (2013)*, pages 119–137. Springer, 2015.
- [24] Jiancheng Weng, Yueyue Wang, Jianling Huang, and Ledian Zhang. Bus operation monitoring oriented public transit travel index system and calculation models. *Advances in Mechanical Engineering*, 2013.
- [25] Hefeng Zhang, Shuxiao Wang, Jiming Hao, Xinming Wang, Shulan Wang, Fahe Chai, and Mei Li. Air pollution and control action in beijing. *Journal of Cleaner Production*, 112:1519–1527, 2016.
- [26] Shaojun Zhang, Ye Wu, Huan Liu, Ruikun Huang, Liuhanzi Yang, Zhenhua Li, Lixin Fu, and Jiming Hao. Real-world fuel consumption and co 2 emissions of urban public buses in beijing. *Applied Energy*, 113:1645–1655, 2014.
- [27] Jiangping Zhou, Enda Murphy, and Ying Long. Commuting efficiency in the beijing metropolitan area: an exploration combining smartcard and travel survey data. *Journal of Transport Geography*, 41:175–183, 2014.